

Operating Systems

1

Significance of this lecture

- This lecture focuses on a node and their central control mechanisms
- How is software linked to physical hardware?
- What controls exist for our physical resources?
- This lesson is concerned with important operating system concepts that impact on distributed system design and behaviour.

2

Topics

- Topics we will cover
 - Operating System Architectures (Monolithic Operating Systems, Microkernels)
 - Kernels
 - Protection
 - Processes and Threads
 - Statefull vs stateless server design
 - Single threaded vs multithreaded server design
 - Architectures for multithreaded servers
 - Thread synchronisation and thread scheduling

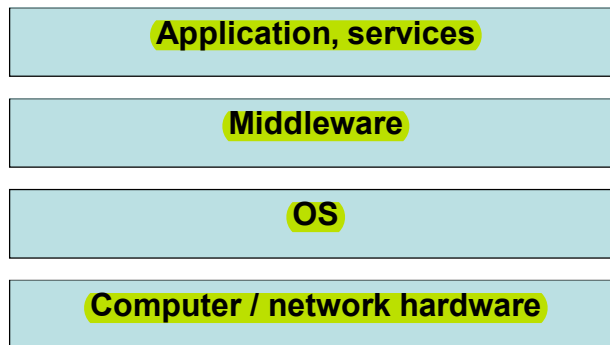
3

The DIS model

- Resource sharing is an important aspect of a DIS
- What is involved in sharing?
 - Clients invoke operations on resources on remote nodes or other local processes
 - This process is simplified through the use of middleware
 - In a DIS client applications/services send request to the middleware
 - The middleware then sends its requests to the host operating system (OS)
 - The OS then utilizes computer and network hardware to send requests
- So “What is the role of the OS?”
 - The OS provide an abstraction of the hardware (network and computer)
 - i.e. processors, memory, communications, storage, etc.
 - Middleware talks to the OS rather than the hardware directly

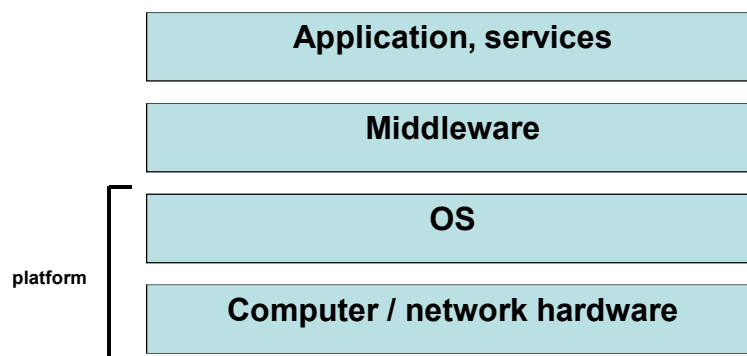
4

Layers of a computer



5

Layers of a computer



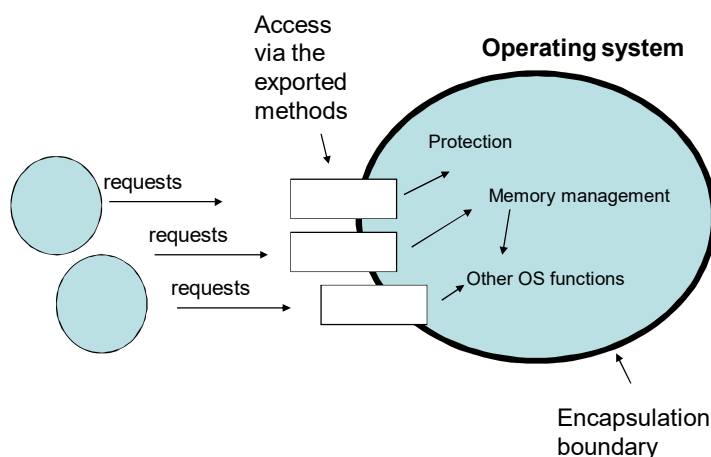
6

OS layers

- Required functionality of an OS
 - **Encapsulation:**
 - With a **useful** set of operations to clients
 - **Unnecessary details** should be transparent (memory management)
 - **Protection:**
 - protection to resources **from incorrect and illegitimate access**
 - **Concurrent processing:**
 - **transparent concurrent process management**
- With respect to remote invocations an OS also needs the following functionality:
 - **Communication:**
 - to remote processes or at least other processes within the computer
 - **Scheduling**
 - Invoked operations need to be schedules to be processed.

7

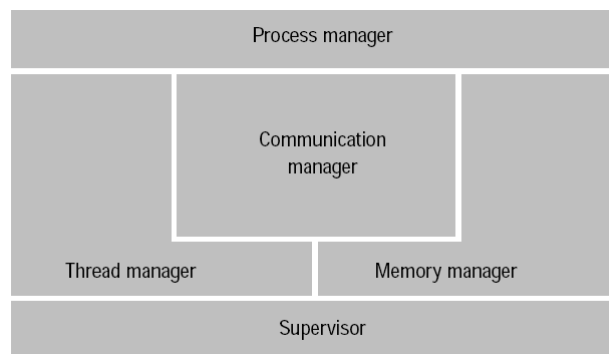
Encapsulation



8

The core OS functions

Micro-kernel architecture



9

OS layers

- **Process manager:**
 - Manages the creation and operations of processes
- **Thread manager:**
 - Creates, synchronizes and schedules threads
- **Communication manager:**
 - Communication between threads on the same host
- **Memory manager:**
 - Physical and virtual memory
- **Supervisor:**
 - The base layer - the hardware abstraction layer
 - Interrupts, system traps, register manipulation, control of the memory management, processor control

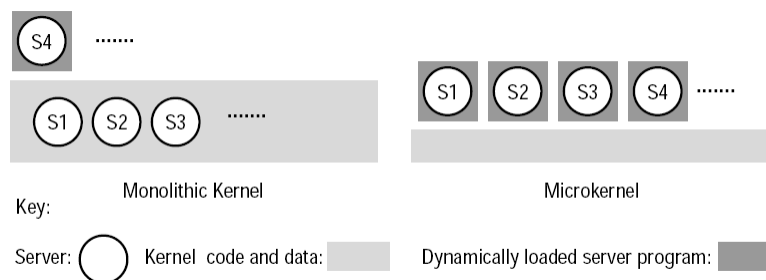
10

Micro/Monolithic kernels

- **Micro kernel**
 - Implement as much as possible outside the kernel
 - Execute less code in Kernel Mode
 - Leads to more context switches
 - Requires inter process messaging
- **Monolithic kernel**
 - Nearly all functions exist within the kernel
 - Executes more code executed in kernel mode
 - Kernels can become very large
 - Faster execution
 - Has high function and code dependence

11

Micro/Monolithic kernels



12

Safeguards and protection

- Resources require protection from illegitimate access
 - illegitimate access is access without privileges
 - An operation is performed that the client does not have access rights to perform
 - illegitimate access is access in unplanned way
 - An operation is performed that is not permitted at all
- Protection is mainly the responsibility of the OS kernel

13

Safeguards and protection

- Features of the OS Kernel:
 - It is always running,
 - Its code has complete access to the physical resources of the computer
 - Processors, memory, etc.
- The kernel allocates resources to processes
 - For example, memory is allocated to user processes and they can only access memory inside its allocated region of memory
- Kernel protection
 - User processes run their own code in their allocated resource space
 - However user processes can also run kernel code
 - They do this through calling kernel supplied handlers
 - These handlers control the access safeguarding against illicit activities
 - Using interrupts or machine-level traps

14

Processes and threads

- A process typically contains an execution environment
 - Execution environment
 - Address space
 - Thread synchronization and communications
 - High level resources (open files etc)
 - Execution environments are expensive to set-up
 - Within its environment a process could decide to run many activities concurrently
 - These concurrent activities are called threads

15

Processes and threads

- Creating a process
 - Creating the process and its execution environment
 - More complex in a DIS
 - Chose a target host
 - Maybe determined by policy
 - i.e. Always at originator
 - Could be based on load balancing (see pp217-218)
 - Create the execution environment on the selected host
 - Two approaches:
 - Share the address space of the existing execution environment
 - Create a new address space

16

Processes and threads

- Threads (Threads of execution)
 - Exist and run within a process execution environment
 - Cannot access threads in other processes
 - Can be dynamically created and destroyed by the process
 - Allow process to perform concurrent operations

17


Processes and threads

- Threads
 - Threads improve performance through concurrency
 - Consider:
 - A single threaded server processes in which it takes 2ms to process a client request
 - Now, assume we have 10 requests = 20ms to process
 - A multi-threaded server processes with 2 threads that can run concurrently
 - We have 10 requests going to 2 different threads
 - 5 requests for each threads so that's 10ms to process
 - With 5 threads we might reduced 4ms
 - Of course there might be other bottlenecks
 - serial hardware such as disk access
 - caching could help for disk access

18

Server threading


- **Thread per request**
 - A new thread is generated for each request
 - They are destroyed when the request completes
 - Increases throughput but introduces creation and destruction overheads
 - This is managed via an I/O thread



19

Server threading


- **Thread per connection**
 - Creates a new thread for each connection that is made to the server
 - Connection could make many requests through its single thread



20

Server threading

- Thread per object
 - Allocates a thread to each remote object
 - This is managed via an I/O thread



21

Client threading

- Client threads
 - Client threading is also beneficial
 - Consider an RMI call that blocks
 - A thread could be generated for the request that would block until its responses
 - Another thread could then continue working whilst the request is serviced

22

Pros and cons of threads

- **Advantages of multi-threads**
 - Threads are more economical to create as they share a execution environment
 - Switching between threads is cheaper than switching between processes
 - Process threads can share data and resources easily
- **Disadvantages of multi-threads**
 - Sharing data and resources circumvents the OS protection
 - A thread might alter data being used by another thread!

23

Thread synchronization

- **Concurrency with threads can lead to errors**
 - as lost updates for example
- **Solution**
 - Use a monitor to specific portions of code, entire methods or objects
 - **A Monitor** is an abstract data type, i.e. a combination of data structures and operations, where at most one thread may be executing at one time.
 - The role of the monitor is to ensure that **only one thread** has access to the specified portion of code at any one time
 - Thus threads might be told to block whilst they wait for access (a notification) that the object or code portion is free

24

Thread scheduling

- Threads need to be told when to **run** and when to **stop**
- There are two main alternatives:
 - **Preemptive scheduling**
 - Greatest control
 - Can have problems with synchronization
 - **Non-preemptive scheduling**
 - Avoids some synchronization problems
 - Offers less control
 - Threads can only be halted when they make a system call
 - Could lead to long executions without an opportunity to halt the process

25

Stateless and Stateful servers

- Information that a server maintains about the status of ongoing interactions with clients is called *state* information
- **Stateless servers**
 - Do not keep state on clients and can change their own state without informing the client
 - Web servers are the classic example of stateless servers
 - Simplifies recovery after crashes
 - Can lead to improved scalability
- **Stateful servers**
 - Maintain state on each of their clients
 - For example file servers are often (but not always) stateful in terms of maintaining information on clients (caching status, etc)
 - Makes crash recovery more complex
 - Can lead to performance optimisation

26

Conclusions

- Operating System Architectures
 - Monolithic kernels
 - Microkernels
- Kernels
- Protection
- Processes and Threads
 - Stateful Vs stateless server design
 - Single threaded Vs multithreaded server design
- Architectures for multithreaded servers
- Thread synchronisation and thread scheduling

27

END

28